

Population Prediction

January 30, 2022

Luis Garduno

Dataset: [International Database \(IDB\)](#)

Question Of Interest: Predict the population of earth in 2122.

1 Data Understanding

1.1 Data Description

```
[1]: import numpy as np
import pandas as pd

# Load dataset into dataframe
df = pd.read_csv('https://raw.githubusercontent.com/luisegarduno/
↳MachineLearning_Projects/master/data/idb5yr.all', delimiter='|',
↳encoding='ISO-8859-1')

df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 34237 entries, 0 to 34236
Data columns (total 99 columns):
#   Column      Non-Null Count  Dtype
---  -
0   #YR          34237 non-null  int64
1   TFR          26163 non-null  float64
2   SRB          26163 non-null  float64
3   RNI          26172 non-null  float64
4   POP95_99     26171 non-null  float64
5   POP90_94     26171 non-null  float64
6   POP85_89     26171 non-null  float64
7   POP80_84     26171 non-null  float64
8   POP75_79     26171 non-null  float64
9   POP70_74     26171 non-null  float64
10  POP65_69     26171 non-null  float64
11  POP60_64     26171 non-null  float64
12  POP5_9       26171 non-null  float64
13  POP55_59     26171 non-null  float64
```

14	POP50_54	26171	non-null	float64
15	POP45_49	26171	non-null	float64
16	POP40_44	26171	non-null	float64
17	POP35_39	26171	non-null	float64
18	POP30_34	26171	non-null	float64
19	POP25_29	26171	non-null	float64
20	POP20_24	26171	non-null	float64
21	POP15_19	26171	non-null	float64
22	POP10_14	26171	non-null	float64
23	POP100_	26171	non-null	float64
24	POPO_4	26171	non-null	float64
25	POP	34237	non-null	int64
26	NMR	26172	non-null	float64
27	NAME	34237	non-null	object
28	MR1_4	26163	non-null	float64
29	MR0_4	26163	non-null	float64
30	MPOP95_99	26171	non-null	float64
31	MPOP90_94	26171	non-null	float64
32	MPOP85_89	26171	non-null	float64
33	MPOP80_84	26171	non-null	float64
34	MPOP75_79	26171	non-null	float64
35	MPOP70_74	26171	non-null	float64
36	MPOP65_69	26171	non-null	float64
37	MPOP60_64	26171	non-null	float64
38	MPOP5_9	26171	non-null	float64
39	MPOP55_59	26171	non-null	float64
40	MPOP50_54	26171	non-null	float64
41	MPOP45_49	26171	non-null	float64
42	MPOP40_44	26171	non-null	float64
43	MPOP35_39	26171	non-null	float64
44	MPOP30_34	26171	non-null	float64
45	MPOP25_29	26171	non-null	float64
46	MPOP20_24	26171	non-null	float64
47	MPOP15_19	26171	non-null	float64
48	MPOP10_14	26171	non-null	float64
49	MPOP100_	26171	non-null	float64
50	MPOPO_4	26171	non-null	float64
51	MPOP	26171	non-null	float64
52	MMR1_4	26163	non-null	float64
53	MMR0_4	26163	non-null	float64
54	IMR_M	26163	non-null	float64
55	IMR_F	26163	non-null	float64
56	IMR	26163	non-null	float64
57	GRR	26163	non-null	float64
58	GR	26172	non-null	float64
59	FPOP95_99	26171	non-null	float64
60	FPOP90_94	26171	non-null	float64
61	FPOP85_89	26171	non-null	float64

```

62 FPOP80_84 26171 non-null float64
63 FPOP75_79 26171 non-null float64
64 FPOP70_74 26171 non-null float64
65 FPOP65_69 26171 non-null float64
66 FPOP60_64 26171 non-null float64
67 FPOP5_9 26171 non-null float64
68 FPOP55_59 26171 non-null float64
69 FPOP50_54 26171 non-null float64
70 FPOP45_49 26171 non-null float64
71 FPOP40_44 26171 non-null float64
72 FPOP35_39 26171 non-null float64
73 FPOP30_34 26171 non-null float64
74 FPOP25_29 26171 non-null float64
75 FPOP20_24 26171 non-null float64
76 FPOP15_19 26171 non-null float64
77 FPOP10_14 26171 non-null float64
78 FPOP100_ 26171 non-null float64
79 FPOPO_4 26171 non-null float64
80 FPOP 26171 non-null float64
81 FMR1_4 26163 non-null float64
82 FMRO_4 26163 non-null float64
83 GENC 34086 non-null object
84 FIPS 34237 non-null object
85 EO_M 26163 non-null float64
86 EO_F 26163 non-null float64
87 EO 26163 non-null float64
88 CDR 26172 non-null float64
89 CBR 26172 non-null float64
90 ASFR45_49 26173 non-null float64
91 ASFR40_44 26173 non-null float64
92 ASFR35_39 26173 non-null float64
93 ASFR30_34 26173 non-null float64
94 ASFR25_29 26173 non-null float64
95 ASFR20_24 26173 non-null float64
96 ASFR15_19 26173 non-null float64
97 AREA_KM2 34237 non-null int64
98 POP_DENS 34237 non-null float64
dtypes: float64(93), int64(3), object(3)
memory usage: 25.9+ MB

```

```

[2]: import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

# Make year column easier to understand
df.rename(columns={'#YR': 'YEAR'}, inplace=True)

```

```
# Remove every column except for year & population
for col in df.columns.values:
    if col != 'YEAR' and col != 'POP':
        df.drop(col, axis=1, inplace=True)

df.describe()
```

```
[2]:
```

	YEAR	POP
count	34237.000000	3.423700e+04
mean	2024.935158	3.306017e+07
std	43.571971	1.253471e+08
min	1950.000000	2.028000e+03
25%	1987.000000	4.304230e+05
50%	2025.000000	4.831844e+06
75%	2063.000000	2.070922e+07
max	2100.000000	1.647894e+09

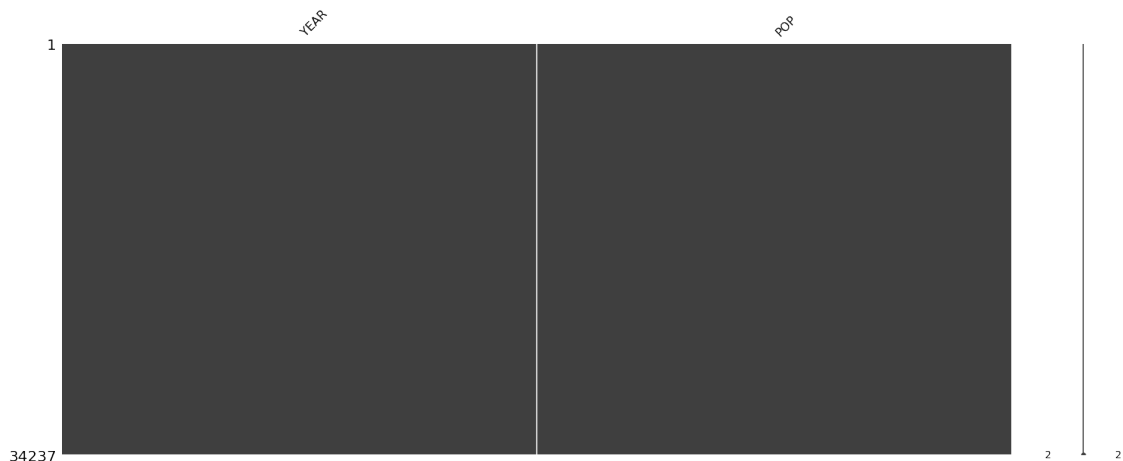
1.2 Data Quality

```
[3]: import missingno as mn

mn.matrix(df)

# Count unique values in column 'gameId' of the dataframe
print('Number of unique values in column "YEAR" : ', df['YEAR'].nunique())
```

Number of unique values in column "YEAR" : 151



1.3 Clearning the Dataset

```
[4]: # Group by year & get sum
df_yr = df.groupby(by='YEAR')
df_yr = df_yr['POP'].sum()

# Create a new dataframe with new data (1951 - 2100)
pop_sum = []
for i in range(1951, 2023):
    pop_sum.append(df_yr[i])
df_pop = pd.DataFrame({'YEAR': list(range(1951, 2023)), 'POP': pop_sum})
df = df_pop


print(f'\n--> Current Population (2021): {df["POP"][70]:,d}\n')
df.tail(5)
```

--> Current Population (2021): 7,831,718,605

```
[4]:      YEAR      POP
67  2018  7597066210
68  2019  7676686052
69  2020  7756873419
70  2021  7831718605
71  2022  7905336896
```

```
[5]: sns.set_style("darkgrid")
plt.subplots(figsize=(10,7))
ax = sns.scatterplot(data=df, x='YEAR', y='POP', color='blue')
ax.set_xlabel('Years', fontsize=16)
ax.set_ylabel('Population', fontsize=16)
ax.set_title('World Population (1951-2022)', fontsize=18)
plt.xlim(1951, 2023)

plt.show()
```



```
[6]: # Define X & Y
if 'POP' in df_pop:
    y = df_pop['POP'].values
    del df_pop['POP']
    X = df_pop.to_numpy()
```

2 Modeling

Derived the formula for calculating the optimal values of the regression weights:

$$w = (X^T X)^{-1} X^T y$$

where X is the matrix of values with a bias column of ones appended onto it. For the population dataset one could construct this X matrix by stacking a column of ones onto the `df_pop.YEAR` matrix.

$$X = \begin{bmatrix} & \vdots & 1 \\ \dots & \text{ds.data} & \dots \\ & \vdots & 1 \end{bmatrix}$$

```
[7]: # Create a matrix full of ones & stack 2 matrices horizontally
X = np.hstack((np.ones((len(X), 1)), X))

# Calculate optimal values of the regression weights
w = np.linalg.inv(X.T @ X) @ X.T @ y

print("\n+++++++ WEIGHTS ++++++\n", pd.DataFrame(data=w))
diff = np.round(( (y - (abs(np.dot(X,w) - y))) / y ) * 100, 2)
print("\n===== TARGET PERCENT ACCURACY =====\n", pd.
↳DataFrame(data=diff))
```

```
+++++++ WEIGHTS ++++++
0
0 -1.494551e+11
1 7.779070e+07

===== TARGET PERCENT ACCURACY =====
0
0 89.20
1 90.73
2 92.10
3 93.32
4 94.38
.. ...
67 99.07
68 99.06
69 99.04
70 99.08
71 99.14

[72 rows x 1 columns]
```

To predict the output from our model, \hat{y} , from w and X we need to use $\hat{y} = w^T X^T$, for row vector \hat{y}

```
[8]: yHat_np = w.T @ X.T          # Shape : (1,72)
yHat_np = yHat_np.ravel()        # Shape : (72,)

MSE_np = (np.square(y - yHat_np)).mean()
```

```
print(f'MSE: {round(MSE_np):,d}')
```

MSE: 7,615,844,338,429,437


```
[9]: X_new = np.array([[0], [2122]])
X_test = np.c_[np.ones((len(X_new), 1)), X_new]
y_test = X_test.dot(w)
print(f'\n--> ~Population (2122): {round(y_test[1]):,d}\n')

sns.set_style("darkgrid")
plt.subplots(figsize=(25,10))

plt.plot(X_new, y_test, "r-", color='red')
sns.scatterplot(data=df, x='YEAR', y=y, color='blue')
sns.scatterplot(x=X_new[1], y=y_test[1], s=100, marker="X", linewidth=1,
               ↪edgecolor='k', color='gold')

plt.xlabel('Years', fontsize=20)
plt.ylabel('Population', fontsize=20)
plt.title('Population Prediction (1951-2122)', fontsize=20)
plt.axis([1951, 2124, 1000000000, 16000000000])
plt.legend(["Linear Regression", "Population", "Prediction"], prop={'size': 15})
plt.show()
```

--> ~Population (2122): 15,616,786,279



3 Comparing Performance

```
[10]: from sklearn.metrics import mean_squared_error
      from sklearn.linear_model import LinearRegression

      reg = LinearRegression().fit(X, y)
      MSE_sk = mean_squared_error(y, reg.predict(X))
      y_testsk = reg.predict(X_test)

      print("***** Linear Equation *****")
      print("[Numpy] \th =", round(w[0],3), "* x + (" + str(round(w[1],5)) + ")")
      print("[Sklearn]\th =", round(reg.intercept_,2), "* x + (" + str(round(reg.
      ↪coef_[1],5)) + ")\n")

      print("***** Mean Squared Error *****")
      print(f'[Numpy] \tMSE: {round(MSE_np):,d}')
      print(f'[Sklearn]\tMSE: {round(MSE_sk):,d}\n')

      print("***** Population Prediction - 2122 *****")
      print(f'[Numpy] \t {round(y_test[1]):,d}')
      print(f'[Sklearn]\t {round(y_testsk[1]):,d}\n')
```

```
***** Linear Equation *****
[Numpy]          h = -149455085136.29 * x + (77790702.83494)
[Sklearn]        h = -149455085136.39 * x + (77790702.83499)

***** Mean Squared Error *****
[Numpy]          MSE: 7,615,844,338,429,437
[Sklearn]        MSE: 7,615,844,338,429,399

***** Population Prediction - 2122 *****
[Numpy]          15,616,786,279
[Sklearn]        15,616,786,279
```

References Census. International Database (IDB). https://www.census.gov/data-tools/demo/idb/#!/country?COUNTRY_YEAR=2022&COUNTRY_YR_ANIM=2022 (Accessed 01-22-2022)